

Linear Predictors

Machine Learning 2023-24

UML Book Chapter 9

Slides P. Zanuttigh

Some material from F. Vandin, N. Ailon, S. Shwartz, J. Janecek

Linear Predictors

The design of a ML strategy requires 2 main steps:

- Select an **hypothesis class \mathcal{H}**
- Select an **algorithm** to find the predictor (i.e. to find $ERM_{\mathcal{H}}(S)$)

For linear models:

Hypotheses Classes

- Halfspaces (binary classification)
- Linear Regression (regression)
- Logistic Regression (classification modeled as a regression problem)

Algorithms

- Linear Programming (for halfspaces, not part of the course)
- Perceptron (for halfspaces)
- Least Squares (for regression)

Affine Function Model

Class of *Affine Functions*:


$$L_d = \{h_{\mathbf{w},b}, \mathbf{w} \in \mathbb{R}^d, b \in \mathbb{R}\}$$

where

$$h_{\mathbf{w},b} = \langle \mathbf{w}, \mathbf{x} \rangle + b = \left(\sum_{i=1}^d w_i x_i \right) + b$$

Each member of L_d is a function $\mathbf{x} \rightarrow \langle \mathbf{w}, \mathbf{x} \rangle + b$, $\mathbf{w} \in \mathbb{R}^d, b \in \mathbb{R}$


- It is a linear function followed by a sum
- Two parameters: b (scalar, called bias) and \mathbf{w} (vector)
- Dimensionality of parameters space: $d+1$



Hypothesis class: $\phi \circ L_d$ $\phi: \mathbb{R} \rightarrow \mathcal{Y}$

- Binary classification $\mathcal{Y} = \{-1, 1\} \rightarrow \phi(z) = \text{sign}(z)$
- Regression $\mathcal{Y} = \mathbb{R} \rightarrow \phi(z) = z$

Geometric Interpretation (2D)



$$x_2 = -\frac{w_1}{w_2}x_1 - \frac{b}{w_2}$$

- The bias is proportional to the offset of the line from the origin
- The weights determine the slope of the line
- The weight vector is perpendicular to the line

Homogeneous Linear Functions

Homogeneous coordinates:

- Idea: incorporate **b** into **w** as an extra dimension/coordinate
- Add an extra dimension to **w**: $\mathbf{w} \rightarrow \mathbf{w}' = \langle b, w_1, w_2, \dots, w_d \rangle$
- Add an extra element to each vector **x**: $\mathbf{x} \rightarrow \mathbf{x}' = \langle 1, x_1, x_2, \dots, x_d \rangle$

Homogeneous linear function:

Rewrite *affine functions*: $L_d = \{h_{\mathbf{w},b}, \mathbf{w} \in \mathbb{R}^d, b \in \mathbb{R}\}$ using homogeneous coord.


$$h_{\mathbf{w},b} = \langle \mathbf{w}, \mathbf{x} \rangle + b = \left(\sum_{i=1}^d w_i x_i \right) + b = b + w_1 x_1 + \dots + w_d x_d$$

$$h_{\mathbf{w}'} = \langle \mathbf{w}', \mathbf{x}' \rangle = \left(\sum_{i=1}^{d+1} w'_i x'_i \right) = b * 1 + w_1 x_1 + \dots + w_d x_d$$

- $\langle \mathbf{w}, \mathbf{x} \rangle + b = \langle \mathbf{w}', \mathbf{x}' \rangle$, rewrite affine function as a linear model
- Get rid of bias (incorporated in the weights vector)
- The affine function becomes a linear function !

Halfspaces Hypothesis Class

Halfspace hypothesis class



- ❑ Input: $\mathcal{X} = \mathbb{R}^d$ (for each sample a vector of features)
 - Using homogeneous coordinates: $\mathbf{x} \rightarrow \mathbf{x}' = (1, x_1, x_2, \dots, x_d) \in \mathbb{R}^{d+1}$
- ❑ Output: $\mathcal{Y} = \{-1, 1\}$ (binary classification)
- ❑ Loss: 0-1 loss

Halfspace Model:

$$HS_d = \text{sign} \circ L_d = \{x \rightarrow \text{sign}(\langle \mathbf{w}, \mathbf{x} \rangle + b), \quad \mathbf{w} \in \mathbb{R}^d, b \in \mathbb{R}\}$$

Using homogeneous coordinates:

$$HS_d = \{x \rightarrow \text{sign}(\langle \mathbf{w}', \mathbf{x}' \rangle), \quad \mathbf{w}' \in \mathbb{R}^{d+1}\}$$


Linear Classification: Halfspace Hypothesis Class

□ $\mathcal{X} = \mathbb{R}^d$, $\mathcal{Y} = \{-1, 1\}$, 0-1 loss

□ Halfspace hypothesis class:

$$HS_d = \text{sign} \circ L_d = \{x \rightarrow \text{sign}(\langle w, x \rangle + b)\}, \quad w \in \mathbb{R}^d, b \in \mathbb{R}$$

Example: $\mathcal{X} = \mathbb{R}^2$

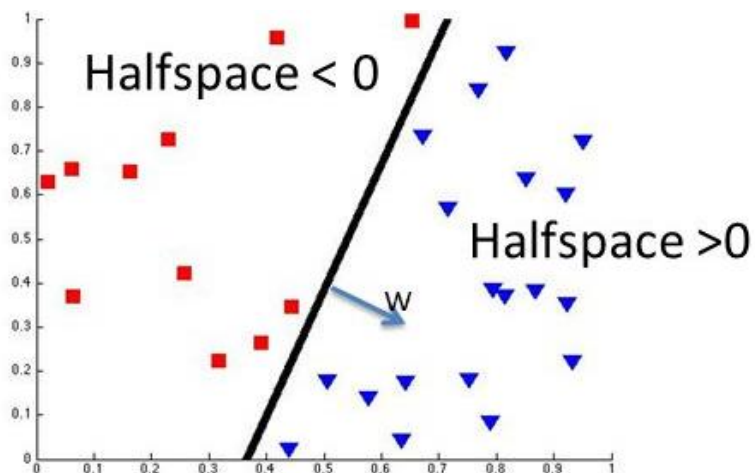


$w \in \mathbb{R}^2, b \in \mathbb{R}$
Using homog. $w' \in \mathbb{R}^3$

Examples: Halfspaces

$$\mathcal{X} = \mathbb{R}^2 \ (d=2)$$


2D space divided by a line



$w \in \mathbb{R}^2, b \in \mathbb{R}$
Using homog. $w' \in \mathbb{R}^3$


$$\mathcal{X} = \mathbb{R}^3 \ (d=3)$$

3D space divided by a plane




$w \in \mathbb{R}^3, b \in \mathbb{R}$
Using homog. $w' \in \mathbb{R}^4$

Realizability: Linearly Separable



Linearly Separable



Not Linearly Separable

Not part of the course

Linear Programming


Linear Program (LP): maximize a **linear** function subject to **linear** inequalities

Target: find $\max_{\mathbf{w} \in \mathbb{R}^d} \langle \mathbf{u}, \mathbf{w} \rangle$ subject to $A\mathbf{w} \geq \mathbf{v}$

$\mathbf{w} \in \mathbb{R}^d$: vector of unknowns, $\mathbf{u} \in \mathbb{R}^d$, $A \in \mathbb{R}^{m \times d}$, $\mathbf{v} \in \mathbb{R}^m$

- ❑ Empirical Risk Minimization (ERM) for halfspaces in the *realizable* case can be expressed as a linear program (LP)
 - The ERM predictor is $L_S(h_S) = 0 \rightarrow \text{sign}(\langle \mathbf{w}, \mathbf{x}_i \rangle) = y_i \forall i \rightarrow y_i \langle \mathbf{w}, \mathbf{x}_i \rangle > 0 \forall i$
 - Need some math to adapt “>” to “≥”
- ❑ All solutions satisfying constraints are ok for us (\rightarrow see SVM later...)
- ❑ There exist efficient LP solvers (e.g., simplex algorithm)

Find ERM Halfspace: Perceptron



- ❑ Iterative algorithm (introduced by Rosenblatt in 1958)
- ❑ Target: find separating hyperplane
- ❑ Find vector \mathbf{w} representing separating hyperplane (in homogeneous coordinates)
- ❑ At each step focus on a misclassified sample and guide the algorithm to be "*more correct*" on it
- ❑ In the realizable case always converge to a (ERM) solution correctly classifying all points
- ❑ Simple and fast in most cases (but there exist critical situations)

Find ERM Halfspace: Perceptron Algorithm

Input: training set $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m)$


initialize $\mathbf{w}^{(1)} = (0, \dots, 0);$

for $t = 1, 2, \dots$ **do**

if $\exists i$ s.t. $y_i \langle \mathbf{w}^{(t)}, \mathbf{x}_i \rangle \leq 0$ **then** $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} + y_i \mathbf{x}_i;$

else return $\mathbf{w}^{(t)};$

Interpretation of update:



Termination? Depends on the realizability assumption!

ERM predictor: $L_S(h_S) = 0$
 $\rightarrow \text{sign}(\langle \mathbf{w}, \mathbf{x}_i \rangle) = y_i \forall i$
 $\rightarrow y_i \langle \mathbf{w}, \mathbf{x}_i \rangle > 0 \forall i$

At each iteration find a misclassified sample and add the sample multiplied by its label

Note that:

$$\begin{aligned} y_i \langle \mathbf{w}^{(t+1)}, \mathbf{x}_i \rangle &= y_i \langle \mathbf{w}^{(t)} + y_i \mathbf{x}_i, \mathbf{x}_i \rangle \\ &= y_i \langle \mathbf{w}^{(t)}, \mathbf{x}_i \rangle + \|\mathbf{x}_i\|^2 \end{aligned}$$

\Rightarrow update guides \mathbf{w} to be “more correct” on (\mathbf{x}_i, y_i) .

$\|\mathbf{x}_i\|^2 > 0$ and target is $y_i \langle \mathbf{w}^{(t)}, \mathbf{x}_i \rangle > 0$: from step t to $t+1$ “more correct” on i -th sample

Perceptron on Linearly Separable Data

Halfspaces:

realizability assumption corresponds to *linearly separable data*

Theorem:

- $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m)$ is linearly separable
- $B = \min\{\|\mathbf{w}\|: y_i \langle \mathbf{w}, \mathbf{x}_i \rangle \geq 1 \quad \forall i = 1, \dots, m\}$
- $R = \max \|\mathbf{x}_i\|$



- The perceptron algorithm stops after at most $(RB)^2$ iterations
- When it stops, it holds that $\forall i \in \{1, \dots, m\}: y_i \langle \mathbf{w}^{(t)}, \mathbf{x}_i \rangle > 0$
- *Notice:* by design the algorithm stops when there are no more wrongly classified samples

Theorem: Demonstration

1. *Define:*

- \mathbf{w}^* : vector achieving the *min* in definition of B
 - Recall $B = \min\{\|\mathbf{w}\| : y_i \langle \mathbf{w}, \mathbf{x}_i \rangle \geq 1 \ \forall i = 1, \dots, m\}$
- T : number of iterations before stopping
 - need to show that $T < (RB)^2$

2. *Consider:* $\frac{\langle \mathbf{w}^*, \mathbf{w}^{(T+1)} \rangle}{\|\mathbf{w}^*\| \|\mathbf{w}^{(T+1)}\|} \rightarrow$ it is smaller than 1 (cosine of angle)

3. We need to demonstrate that:

$$\frac{\sqrt{T}}{RB} = \frac{T}{\sqrt{T}RB} \leq \frac{\langle \mathbf{w}^*, \mathbf{w}^{(T+1)} \rangle}{\|\mathbf{w}^*\| \|\mathbf{w}^{(T+1)}\|} \leq 1 \Rightarrow T < (RB)^2$$

4. Proceed in 2 parts:

- a) Numerator: demonstrate that $\langle \mathbf{w}^*, \mathbf{w}^{(T+1)} \rangle \geq T$
- b) Denominator: demonstrate that $\|\mathbf{w}^*\| \|\mathbf{w}^{(T+1)}\| \leq \sqrt{T} RB$

Theorem: Demonstration (Numerator)

Numerator: demonstrate that $\langle \mathbf{w}^*, \mathbf{w}^{(T+1)} \rangle \geq T$

- First iteration : $\mathbf{w}^{(1)} = (0, \dots, 0) \rightarrow \langle \mathbf{w}^*, \mathbf{w}^{(1)} \rangle = 0$
- At each step $\langle \mathbf{w}^*, \mathbf{w}^{(t+1)} \rangle - \langle \mathbf{w}^*, \mathbf{w}^{(t)} \rangle \geq 1$ (using perceptron update rule and recalling definition of \mathbf{w}^* , see *)
- After T iterations: $\langle \mathbf{w}^*, \mathbf{w}^{(T+1)} \rangle \geq T$ (see *)

$$\begin{aligned}
 (*) \quad \langle \mathbf{w}^*, \mathbf{w}^{(T+1)} \rangle &= \sum_{t=1}^T (\langle \mathbf{w}^*, \mathbf{w}^{(t+1)} \rangle - \langle \mathbf{w}^*, \mathbf{w}^{(t)} \rangle) \\
 &= \sum_{t=1}^T \langle \mathbf{w}^*, \mathbf{w}^{(t+1)} - \mathbf{w}^{(t)} \rangle = \sum_{t=1}^T \langle \mathbf{w}^*, y_i \mathbf{x}_i \rangle \geq T
 \end{aligned}$$

Algorithm
Assumption on \mathbf{w}^*

(perceptron update rule)
(definition of \mathbf{w}^* and B)
 $\rightarrow \langle \mathbf{w}^*, y_i \mathbf{x}_i \rangle \geq 1$

Theorem: Demonstration (Denominator)

Denominator: demonstrate that $\|\mathbf{w}^*\| \|\mathbf{w}^{(T+1)}\| \leq \sqrt{T} RB$

a) $\|\mathbf{w}^{(T+1)}\|^2 \leq TR^2 \rightarrow \|\mathbf{w}^{(T+1)}\| \leq \sqrt{T}R \quad (**)$

b) $\|\mathbf{w}^*\| = B$ (by definition)

$$\begin{aligned}
 (**) \quad \|\mathbf{w}^{(T+1)}\|^2 &= \sum_{t=1}^T \left(\|\mathbf{w}^{(t+1)}\|^2 - \|\mathbf{w}^{(t)}\|^2 \right) \\
 &= \sum_{t=1}^T \left(\|\mathbf{w}^{(t)} + y_i \mathbf{x}_i\|^2 - \|\mathbf{w}^{(t)}\|^2 \right) \quad \leftarrow \text{the two } \|\mathbf{w}^{(t)}\|^2 \text{ cancel out} \\
 &= \sum_{t=1}^T \left(2y_i \langle \mathbf{w}^{(t)}, \mathbf{x}_i \rangle + \|\mathbf{x}_i\|^2 \right) \leq TR^2 \\
 &\quad \leq 0 \text{ by algorithm} \quad \leq R^2 \text{ by algorithm} \\
 &\quad \text{(perceptron update condition: select misclassified sample)} \quad \text{(definition of R)} \\
 &\quad \quad \quad R = \max \|\mathbf{x}_i\|
 \end{aligned}$$

Perceptron: Notes

It is simple to implement !

On separable data:

- ☐ Convergence is guaranteed
- ☐ Convergence depends on B , which can be exponential in d
 - If the input vectors are not normalized and arranged in some unfavorable ways the running time can be very long
 - A Linear Programming (LP) approach may be better to find ERM solution in some cases
- ☐ Potentially multiple solutions, which one is picked depends on starting values

On non separable data:

- ☐ Run for some time and keep best solution found up to that point (*pocket algorithm*)

Learning Example


Initial Values:

$$\eta = 0.2$$

Learning rate
(also denoted with λ)

$$w = \begin{pmatrix} 0 \\ 1 \\ 0.5 \end{pmatrix}$$

Initial line
position



$$\begin{aligned} 0 &= w_0 + w_1 x_1 + w_2 x_2 \\ &= 0 + x_1 + 0.5 x_2 \\ \Rightarrow x_2 &= -2x_1 \end{aligned}$$

Perceptron with learning rate: $w^{(t+1)} = w^{(t)} + \lambda y_i x_i$

Learning Example


$$\eta = 0.2$$

$$w = \begin{pmatrix} 0 \\ 1 \\ 0.5 \end{pmatrix}$$

$$x_1 = 1, x_2 = 1$$

$$w^T x > 0$$

Correct classification,
no action



Learning Example

$$\eta = 0.2$$

$$\mathbf{w} = \begin{pmatrix} 0 \\ 1 \\ 0.5 \end{pmatrix}$$


$$x_1 = 2, x_2 = -2$$

$$w_0 = w_0 - 0.2 * 1$$

$$w_1 = w_1 - 0.2 * 2$$

$$w_2 = w_2 - 0.2 * (-2)$$

$$\text{new } \mathbf{w} = \begin{bmatrix} -0.2 \\ 0.6 \\ 0.9 \end{bmatrix}$$




Learning Example

$$\eta = 0.2$$

$$w = \begin{pmatrix} -0.2 \\ 0.6 \\ 0.9 \end{pmatrix}$$

previous w

$$w = \begin{pmatrix} 0 \\ 1 \\ 0.5 \end{pmatrix}$$



$$x_1 = 2, x_2 = -2$$

$$w_0 = w_0 - 0.2 * 1$$

$$w_1 = w_1 - 0.2 * 2$$

$$w_2 = w_2 - 0.2 * (-2)$$

Learning Example


$$\eta = 0.2$$

$$w = \begin{pmatrix} -0.2 \\ 0.6 \\ 0.9 \end{pmatrix}$$

$$x_1 = -1, x_2 = -1.5$$

$$w^T x < 0$$

Correct classification
no action



Learning Example


$$\eta = 0.2$$

$$w = \begin{pmatrix} -0.2 \\ 0.6 \\ 0.9 \end{pmatrix}$$

$$x_1 = -2, x_2 = -1$$

$$w^T x < 0$$

Correct classification
no action



Learning Example

$$\eta = 0.2$$

$$w = \begin{pmatrix} -0.2 \\ 0.6 \\ 0.9 \end{pmatrix}$$


$$x_1 = -2, x_2 = 1$$

$$w_0 = w_0 + 0.2 * 1$$

$$w_1 = w_1 + 0.2 * (-2)$$

$$w_2 = w_2 + 0.2 * 1$$

$$\text{new } w = \begin{bmatrix} 0 \\ 0.2 \\ 1.1 \end{bmatrix}$$




Learning Example

$$\eta = 0.2$$

$$w = \begin{pmatrix} 0 \\ 0.2 \\ 1.1 \end{pmatrix}$$

previous w

$$w = \begin{pmatrix} -0.2 \\ 0.6 \\ 0.9 \end{pmatrix}$$



$$x_1 = -2, x_2 = 1$$

$$w_0 = w_0 + 0.2 * 1$$

$$w_1 = w_1 + 0.2 * (-2)$$

$$w_2 = w_2 + 0.2 * 1$$

Learning Example

$$\eta = 0.2$$

$$\mathbf{w} = \begin{pmatrix} 0 \\ 0.2 \\ 1.1 \end{pmatrix}$$


$$x_1 = 1.5, x_2 = -0.5$$

$$w_0 = w_0 + 0.2 * 1$$

$$w_1 = w_1 + 0.2 * 1.5$$

$$w_2 = w_2 + 0.2 * (-0.5)$$

$$\text{new } \mathbf{w} = \begin{bmatrix} 0.2 \\ 0.5 \\ 1 \end{bmatrix}$$



Learning Example

$$\eta = 0.2$$

$$w = \begin{pmatrix} 0.2 \\ 0.5 \\ 1 \end{pmatrix}$$

previous w


$$w = \begin{pmatrix} 0 \\ 0.2 \\ 1.1 \end{pmatrix}$$

$$x_1 = 1.5, x_2 = -0.5$$

$$w_0 = w_0 + 0.2 * 1$$

$$w_1 = w_1 + 0.2 * 1.5$$

$$w_2 = w_2 + 0.2 * (-0.5)$$



All samples correctly classified \rightarrow perceptron algorithm stops !

VC Dimension of Halfspaces


For the halfspace hypothesis class:

- The VC dimension of the class of **homogenous** halfspaces in \mathbb{R}^d is d
- The VC dimension of the class of **non-homogenous** halfspaces in \mathbb{R}^d is $d+1$


Example: 2D space

- *i.e., the hyperplane is a line*
 $d=2$ in non-homogeneous or $d+1=3$ in homogeneous coord.
- $VCdim(\mathcal{H}) \geq 3$ (see example)
- $VCdim(\mathcal{H}) < 4$ (no set of size 4 can be shattered)

$VCdim(\mathcal{H}) \geq 3$



$VCdim(\mathcal{H}) < 4$



VC Dimension of Halfspaces: Demonstration (1)

The VC dimension of the class of **homogenous** halfspaces in \mathbb{R}^d is d

Demonstration (homogenous case):

$$a) VCdim(HS_d^{hmg}) \geq d$$

1. Consider the set $\mathbf{e}_1, \dots, \mathbf{e}_d$ where $\forall i: \mathbf{e}_i = (0, \dots, 0, 1, 0, \dots, 0)$
 - i.e., all "0" except "1" in the i -th coordinate
2. The set is shattered by the homogeneous halfspace: to obtain the labeling $\mathbf{e}_1, \dots, \mathbf{e}_d$ set $\mathbf{w} = (y_1, \dots, y_d) \Rightarrow \langle \mathbf{w}, \mathbf{e}_i \rangle = y_i \quad \forall i$
 - for each vector only the multiplication with the corresponding label is $\neq 0$ (only the i -th term remains)
$$\langle (y_1, \dots, y_{i-1}, y_i, y_{i+1}, \dots, y_d), (0, \dots, 0, 1, 0, \dots, 0) \rangle = y_i$$

VC Dimension of Halfspaces: Demonstration (2)

$$b) VCdim(HS_d^{hmg}) < d + 1$$

1. $\mathbf{x}_1, \dots, \mathbf{x}_{d+1}$ generic set of $d+1$ vectors in \mathbb{R}^d

2. They must be linearly dependent:

$$\exists a_1, \dots, a_{d+1} \in \mathbb{R} \text{ (not all zero): } \sum_{i=1}^{d+1} a_i \mathbf{x}_i = 0$$

3. Define $I = \{i: a_i > 0\}$, $J = \{j: a_j < 0\}$: either I or J are non-empty

4. Assume both non-empty: $\sum_{i \in I} a_i \mathbf{x}_i = \sum_{j \in J} |a_j| \mathbf{x}_j$

5. *By contradiction*: assume that the set is shattered: \exists a vector \mathbf{w} such that $\langle \mathbf{w}, \mathbf{x}_i \rangle > 0 \forall i \in I$ and $\langle \mathbf{w}, \mathbf{x}_j \rangle < 0 \forall j \in J$

6. It follows a contradiction :

$$0 < \sum_{i \in I} a_i \langle \mathbf{x}_i, \mathbf{w} \rangle = \langle \sum_{i \in I} a_i \mathbf{x}_i, \mathbf{w} \rangle = \langle \sum_{j \in J} |a_j| \mathbf{x}_j, \mathbf{w} \rangle = \sum_{j \in J} |a_j| \langle \mathbf{x}_j, \mathbf{w} \rangle < 0$$

7. If I or J are empty just replace one of the two inequalities with "=" but still there is the contradiction!!

Linear Regression

Regression: estimate the relation between some explanatory variables (*features*) and some real valued outcome

- ❑ Domain set : $\mathcal{X} \in \mathbb{R}^d$, label set : $\mathcal{Y} = \mathbb{R}$
- ❑ Find $h \in \mathcal{H}_{reg}: \mathbb{R}^d \rightarrow \mathbb{R}$ that best approximates the relation between input and output
- ❑ Hypothesis class (*linear regression*):

$$\mathcal{H}_{reg} = L_d = \{\mathbf{x} \rightarrow \langle \mathbf{w}, \mathbf{x} \rangle + b : \mathbf{w} \in \mathbb{R}^d, b \in \mathbb{R}\}$$

- ❑ Loss function: *squared loss (L2, MSE)* is commonly used but other functions are possible (e.g., *mean absolute error*)


$$\ell(h, (\mathbf{x}, y)) \stackrel{\text{def}}{=} (h(\mathbf{x}) - y)^2$$

- ❑ Empirical Risk function: *Mean Squared Error* on training set

$$L_s(h) = \frac{1}{m} \sum_{i=1}^m (h(\mathbf{x}_i) - y_i)^2$$



Linear Regression (1D)

$$\mathcal{X} = \mathbb{R}^1 \quad \mathcal{Y} = \mathbb{R}$$




Example: Linear Regression (1)

Online Store	Monthly Sales (in 1000 \$)	Online Advertising Dollars (1000 \$)
1	368	1.7
2	340	1.5
3	665	2.8
4	954	5.0
5	331	1.3
6	556	2.2
7	376	1.3




Example: Linear Regression (2)

Car Age (years)	Price (€)
4	6300
4	5800
5	5700
5	4500
7	4500
7	4200
8	4100
9	3100
10	2100
11	2500
12	2200



Example: Linear Regression (3)

- $\mathcal{X} \subset \mathbb{R}^d$, $\mathcal{Y} \subset \mathbb{R}$, $\mathcal{H} = \{\mathbf{x} \mapsto \langle \mathbf{w}, \mathbf{x} \rangle : \mathbf{w} \in \mathbb{R}^d\}$
- Example: $d = 1$, predict weight of a child based on his age.



Least Squares

$$\arg \min_{\mathbf{w}} L_s(h_{\mathbf{w}}) = \arg \min_{\mathbf{w}} \frac{1}{m} \sum_{i=1}^m (\langle \mathbf{w}, \mathbf{x}_i \rangle - y_i)^2$$

- ❑ The *least squares* algorithm solves the ERM problem for *linear regression* predictors with the *squared loss*
- ❑ Find the parameters vector that minimize the MSE between the estimated and training values
- ❑ To solve the problem: calculate gradient *w.r.t vector \mathbf{w}* and set to 0

Least Squares: Solution

$$\arg \min_{\mathbf{w}} L_s(h_{\mathbf{w}}) = \arg \min_{\mathbf{w}} \frac{1}{m} \sum_{i=1}^m (\langle \mathbf{w}, \mathbf{x}_i \rangle - y_i)^2$$

- Compute gradient w.r.t \mathbf{w} and set to 0

$$\frac{\partial L_s}{\partial \mathbf{w}} = \frac{2}{m} \sum_{i=1}^m (\langle \mathbf{w}, \mathbf{x}_i \rangle - y_i) \mathbf{x}_i = 0 \rightarrow \sum_{i=1}^m \langle \mathbf{w}, \mathbf{x}_i \rangle \mathbf{x}_i = \sum_{i=1}^m y_i \mathbf{x}_i$$

- Set

$$A = \left(\sum_{i=1}^m \mathbf{x}_i \mathbf{x}_i^T \right) = \begin{bmatrix} \vdots & & \\ \mathbf{x}_1 & & \\ \vdots & & \end{bmatrix} \dots \begin{bmatrix} \vdots & & \\ \mathbf{x}_m & & \\ \vdots & & \end{bmatrix} \begin{bmatrix} \dots & \mathbf{x}_1 & \dots \\ \vdots & & \\ \dots & \mathbf{x}_m & \dots \end{bmatrix} \quad \mathbf{b} = \sum_{i=1}^m y_i \mathbf{x}_i = \begin{bmatrix} \vdots & & \vdots \\ \mathbf{x}_1 & \dots & \mathbf{x}_m \\ \vdots & & \vdots \end{bmatrix} \begin{bmatrix} y_1 \\ \vdots \\ y_m \end{bmatrix}$$

- The solution is:

$$\sum_{i=1}^m \langle \mathbf{w}, \mathbf{x}_i \rangle \mathbf{x}_i = \sum_{i=1}^m y_i \mathbf{x}_i \rightarrow A\mathbf{w} = \mathbf{b} \rightarrow \mathbf{w} = A^{-1}\mathbf{b}$$

- The unknown is \mathbf{w} , A : $d \times d$ matrix, \mathbf{b} and \mathbf{w} : d -dimensional vectors
- Works also in the case in which A is not invertible, but this requires a special handling (*not part of the course*)

Polynomial Regression

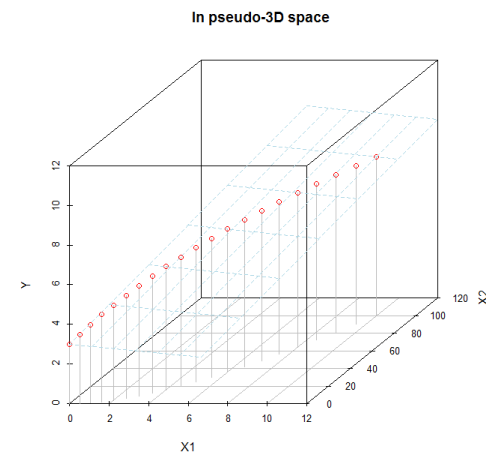

- Polynomial regression: find the one dimensional polynomial of degree n that better predicts the data
 - $p(x) = a_0 + a_1x + a_2x^2 + \dots + a_nx^n$
 - Need to estimate the coefficient vector \mathbf{a}
 - 1D polynomial pred. deg. n : $\mathcal{H}_{poly}^n = \{x \rightarrow p(x)\}, \mathcal{X} = \mathbb{R}, \mathcal{Y} = \mathbb{R}$
- Reduce the problem to a n -dimensional linear regression using the mapping:

$$\psi: \mathbb{R} \rightarrow \mathbb{R}^{n+1} \quad \psi(x) = (1, x, x^2, \dots, x^n)$$

- We obtain:


$$\langle \mathbf{a}, \psi(x) \rangle = a_0 + a_1x + a_2x^2 + \dots + a_nx^n$$

- Find the vector of coefficients \mathbf{a} using the Least Square algorithm
- Non-linear relation becomes linear in the higher dimensional space
- Notice that the variables are not independent
- The optimization can become unstable for large n




Logistic Regression


- ❑ Reframe a classification problem as a regression one
- ❑ Target as in regression:
 - learn a function $h: \mathbb{R}^d \rightarrow [0; 1]$
 - the output of h is a real number
- ❑ Used for classification:
 - interpret the output of h as the probability that the label is 1
 - *regression-like output for classification !*
- ❑ We'll deal with binary classification but the approach can be extended to the multi-class setting
- ❑ Hypothesis class $\mathcal{H}: \phi_{sig} \circ L_d$ where $\phi: \mathbb{R} \rightarrow [0,1]$ is the sigmoid function and L_d a linear function



Sigmoid Function



$$\phi_{sig}(z) = \frac{1}{1 + e^{-z}}$$



- ❑ Bigger than $\frac{1}{2}$ for positive values and smaller for negative ones
- ❑ Tends to 1 for large positive values and to 0 for negative ones
- ❑ Can be viewed as a scaled and shifted "soft" sign function

Loss for Logistic Regression

- Instead of hard choice \rightarrow use probability of correct label being 0 or 1

$$H_{sig} = \phi_{sig} \circ L_d = \{x \rightarrow \phi_{sig}(\langle w, x \rangle) : w \in \mathbb{R}^d\}$$

- Hypothesis class : $h_w(x) = \frac{1}{1+e^{-\langle w, x \rangle}}$

- Loss function: $\ell(h_w, (x, y)) = \log(1 + e^{-y\langle w, x \rangle})$

- ERM Problem: $\operatorname{argmin}_{w \in \mathbb{R}^d} \frac{1}{m} \sum_{i=1}^m \log(1 + e^{-y_i \langle w, x_i \rangle})$

Logistic Loss Function

Loss function: $\ell(h_w, (x, y)) = \log(1 + e^{-y\langle w, x \rangle})$, why ?

Consider $h_w(x) = \frac{1}{1 + e^{-\langle w, x \rangle}} \leftrightarrow y \in \{+1, -1\}$

□ Case $y=1$: need $h_w(x) \rightarrow 1$

$$h_w(x) = \frac{1}{1 + e^{-\langle w, x \rangle}} = \frac{1}{1 + e^{-y\langle w, x \rangle}}$$

- If denominator small $h_w(x) \rightarrow 1$ good case
- If denominator large $h_w(x) \rightarrow 0$ error

□ Case $y=-1$: need $h_w(x) \rightarrow 0 \Rightarrow 1 - h_w(x) \rightarrow 1$

$$1 - h_w(x) = 1 - \frac{1}{1 + e^{-\langle w, x \rangle}} = \frac{1 + e^{-\langle w, x \rangle} - 1}{1 + e^{-\langle w, x \rangle}} = \frac{1}{e^{\langle w, x \rangle} + 1} = \frac{1}{1 + e^{-y\langle w, x \rangle}}$$

➤ Same as before :

- If denominator small $1 - h_w(x) \rightarrow 1$ good case
- if denominator large $1 - h_w(x) \rightarrow 0$ error

□ Loss need to increase with $1 + e^{-y\langle w, x \rangle}$ and log function is monotonic

Maximum Likelihood Estimation (MLE)

Not part of the course

Maximum Likelihood Estimation (MLE) is a statistical approach for finding the parameters that maximize the joint probability of a given dataset assuming a specific parametric probability function

- ❑ MLE essentially assumes a generative model for the data
- ❑ MLE solution is equivalent to ERM solution for logistic regression

MLE approach:

1. Given training set $S = ((\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m))$, assume each (\mathbf{x}_i, y_i) is i.i.d. from some probability distribution (that is characterized by some parameters)
2. Consider $P[S|\theta]$ (likelihood of data given parameters)
3. log likelihood: $L(S; \theta) = \log(P[S|\theta])$
 - *log*: monotonic \rightarrow same maximum, but simpler to differentiate
4. *Maximum Likelihood Estimator (MLE)*: $\hat{\theta} = \operatorname{argmax}_{\theta} L(S; \theta)$

MLE and Logistic Regression

Not part of the course

MLE solution is equivalent to ERM solution for logistic regression

Logistic Regression:

1. Assume training set $S = ((x_1, y_1), \dots, (x_m, y_m))$
2. $P[y_i = 1] = h_{\mathbf{w}}(x_i) = \frac{1}{1+e^{-\langle \mathbf{w}, x_i \rangle}} = \frac{1}{1+e^{-y_i \langle \mathbf{w}, x_i \rangle}}$ (since $y_i = 1$)
3. $P[y_i = -1] = 1 - h_{\mathbf{w}}(x_i) = \frac{1}{1+e^{\langle \mathbf{w}, x_i \rangle}} = \frac{1}{1+e^{-y_i \langle \mathbf{w}, x_i \rangle}}$ (first equality recall logistic loss, 2nd since $y_i = -1$)
4. Likelihood of training set (joint probability $P[S|\mathbf{w}]$) : $\prod_{i=1}^m \left(\frac{1}{1+e^{-y_i \langle \mathbf{w}, x_i \rangle}} \right)$
5. Log likelihood : $\log(P[S|\mathbf{w}]) = \log \prod_{i=1}^m \left(\frac{1}{1+e^{-y_i \langle \mathbf{w}, x_i \rangle}} \right) = -\sum_{i=1}^m \log(1 + e^{-y_i \langle \mathbf{w}, x_i \rangle})$
*→ corresponds to (-1)*logistic loss*

Maximum Likelihood Estimator:



$$\operatorname{argmax}_{\mathbf{w}} L(S; \mathbf{w}) = \operatorname{argmax}_{\mathbf{w}} \log(P[S|\mathbf{w}]) = \operatorname{argmin}_{\mathbf{w}} \sum_{i=1}^m \log(1 + e^{-y_i \langle \mathbf{w}, x_i \rangle})$$

Recall: $\operatorname{argmax}(-x) = \operatorname{argmin}(x)$ They have the same target !!!

Not part of the course

DELL'INFORMAZIONE

Example: MLE for Gaussian PDF (1)



Assume that the data is produced by a Gaussian distribution

$$P(x; \mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

$$f_1 \sim N(10, 2.25) \quad f_2 \sim N(10, 9), \\ f_3 \sim N(10, 0.25) \quad f_4 \sim N(8, 2.25)$$

find μ, σ maximizing the joint probability of data

Example: MLE for Gaussian PDF (2)

Not part of the course

Joint probability

$$P(9, 9.5, 11; \mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(9-\mu)^2}{2\sigma^2}\right) \times \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(9.5-\mu)^2}{2\sigma^2}\right) \times \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(11-\mu)^2}{2\sigma^2}\right)$$

Assume
3 samples:
9, 9.5, 11

Log likelihood

$$\ln(P(x; \mu, \sigma)) = \ln\left(\frac{1}{\sigma\sqrt{2\pi}}\right) - \frac{(9-\mu)^2}{2\sigma^2} + \ln\left(\frac{1}{\sigma\sqrt{2\pi}}\right) - \frac{(9.5-\mu)^2}{2\sigma^2} + \ln\left(\frac{1}{\sigma\sqrt{2\pi}}\right) - \frac{(11-\mu)^2}{2\sigma^2}$$

Recall
 $\log(ab) = \log(a) + \log(b)$

Differentiate w.r.t μ and set to 0

Get optimal mean

$$\frac{\partial \ln(P(x; \mu, \sigma))}{\partial \mu} = \frac{1}{\sigma^2} [9 + 9.5 + 11 - 3\mu] = 0$$

The same can be done for σ

set derivative to 0

$$\mu = \frac{9 + 9.5 + 11}{3} = 9.833$$